

Big Data: Big N or Very Large Sample Case

VC

May 13, 2016

Examples of Very Big Data

- ▶ Congressional record text, in 100 GBs
- ▶ Nielsen's scanner data, 5TBs
- ▶ Medicare claims data are in 100 TBs
- ▶ Facebook 200,000 TBs

Map Reduce & Hadoop

The basic idea is that you need to divide work among the cluster of computers since you can't store and analyze the data on a single computer.

Simple but powerful algorithm framework. Released by Google around 2004; Hadoop is an open-source version.

Map-Reduce algorithm has the following steps:

1. Map: processes "chunks" of data to produce "summaries"
2. Reduce: combines "summaries" from different chunks to produce a single output file

Formally, if data $D = \cup_{i=1}^K D_i$,

$$f(D) = R(M(D_1), \dots, M(D_K)),$$

where

- ▶ $M(\cdot)$ is the mapper,
- ▶ $R(\cdot)$ is the reducer.

Examples

- ▶ Count words in doc i , D_i . Map: $M : D_i \mapsto M(D_i)$, set of (*word*, *count*) pairs, also called "key-value" pairs. Reduce: Aggregate $\{M(D_i)\}_{i=1}^K$ by summing over *count* within *word*.
- ▶ Records at hospital i , D_i . Map: $M : D_i \mapsto M(D_i)$, say records for patients who are 65+. Reduce: take union over elements of $\{M(D_i)\}$, i.e.

$$f(D) = R(\{M(D_i)\}_{i=1}^k) = \cup_{i=1}^k \{M(D_i)\}.$$

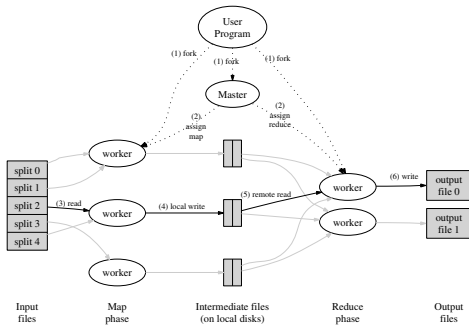
- ▶ Compute the minimum of data $D = \cup_{i=1}^k D_i$, where D_i are vectors. Map: $M : D_i \mapsto M(D_i) = \min(D_i)$. Reduce: take minimum over elements of $\{M(D_i)\}$, i.e.

$$f(D) = R(\{M(D_i)\}_{i=1}^k) = \min_{i=1}^k \{\min(D_i)\}.$$

Map-Reduce Functionality

- ▶ Partitions data across machines
- ▶ Schedules execution across nodes
- ▶ Manages communication across machines
- ▶ Handles errors, machine failure

MapReduce: Model



Amazon Web Services

- ▶ Data centers owned and run by Amazon. You can rent "virtual computers" minute-by-minute basis
- ▶ more than 80% of the cloud computing market
- ▶ nearly 3,000 employees
- ▶ cost per machine: 0.01 to 4.00 /hour
- ▶ Several services in AWS
- ▶ S3 (Storage)
- ▶ EC2 (Individual Machines)
- ▶ Elastic Map Reduce
- ▶ distribute the data for Hadoop clusters

Distributed and Recursive Computing of Estimators

We want to compute the least squares estimator

$$\hat{\beta} \in \arg \min_b n^{-1} \sum_{i=1}^n (y_i - x_i' b)^2.$$

The sample size n is very large and can't load the data into a single machine. What could we do if we have a single machine or many machines?

Use the classical sufficiency ideas to distribute jobs across machines, spatially or in time.

The OLS Example

- ▶ We know that

$$\hat{\beta} = (X'X)^{-1}(X'Y).$$

- ▶ Hence we can do everything we want with just:

$$X'X, \quad X'Y, \quad n, \quad S_0,$$

where S_0 is a "small" random sample $(Y_i, X_i)_{i \in I_0}$ with sample size n_0 , where n_0 is large, but small enough that the data can be loaded in the machine.

- ▶ We need $X'X$ and $X'Y$ to compute the estimators to compute the estimator.
- ▶ We need S_0 to compute robust standard errors and we need to know n to scale these standard errors appropriately.

The OLS Example Continued

- ▶ The terms like $X'X$ and $X'Y$ are sums that can be computed by distribution of jobs over many machines:
1. Suppose machine j stores sample $S_j = (X_i, Y_i)_{i \in I_j}$ of size n_j .
 2. Then we can map S_j to the sufficient statistics

$$T_j = \left(\sum_{i \in I_j} X_i X_i', \sum_{i \in I_j} X_i Y_i, n_j \right)$$

for each j .

3. We then collect $(T_j)_{j=1}^M$ and reduce them further to

$$T = \sum_{j=1}^M T_j = (X'X, X'Y, n).$$

The LASSO Example

The Lasso estimator minimizes

$$(Y - X\beta)'(Y - X\beta) + \lambda\|\Psi\beta\|_1, \quad \Psi = \text{diag}(X'X)$$

or equivalently

$$Y'Y - 2\beta'X'Y + \beta'X'X\beta + \lambda\|\Psi\beta\|_1.$$

Hence in order to compute Lasso and estimate noise level to tune λ we only need to know

$$Y'X, \quad X'X, \quad n, \quad S_0.$$

Computation of sums could be distributed across machines.

The Two Stage Least Squares

The estimator takes the form

$$(X'P_ZX)^{-1}X'P_ZY = (X'Z(Z'Z)^{-1}Z'X)^{-1}X'Z(Z'Z)^{-1}Z'Y.$$

Thus we only need to know

$$Z'Z, \quad X'Z, \quad Z'Y, \quad n, \quad S_0.$$

Computation of sums could be distributed across machines.

Exponential Families and Non-Linear Examples

Consider estimation using MLE based upon exponential families. Here assume data $W_i \sim f_\theta$, where

$$f_\theta(w) = \exp(T(w)' \theta + \varphi(\theta)).$$

Then the MLE maximizes

$$\sum_{i=1}^n \log f_\theta(W_i) = \sum_{i=1}^n T(W_i)' \theta + \varphi(\theta) =: T' \theta + n \varphi(\theta).$$

The sufficient statistic T can be obtained via distributed computing. We also need an S_0 to obtain standard errors. Going beyond such quasi-linear examples could be difficult, but possible.

M- and GMM - Estimation

The ideas could be pushed forward using 1-step or approximate minimization principles. Here is a very crude form of one possible approach.

Suppose that $\hat{\theta}$ minimizes

$$\sum_{i=1}^n m(W_i, \theta).$$

Then given an initial estimator $\hat{\theta}_0$ computed on S_0 we could do Gradient descent iterations to approximate $\hat{\theta}$:

$$\hat{\theta}_{j+1} = \hat{\theta}_j - A_j \sum_{i=1}^n \nabla_{\theta} m(W_i, \hat{\theta}_j).$$

where A_j are matrices, e.g. $A_j = a_j I$.

We could also do Newton iterations to approximate $\hat{\theta}$:

$$\hat{\theta}_{j+1} = \hat{\theta}_j - \left(\sum_{i=1}^n \nabla_{\theta}^2 m(W_i, \hat{\theta}_j) \right)^{-1} \sum_{i=1}^n \nabla_{\theta} m(W_i, \hat{\theta}_j).$$

M- and GMM - Estimation

Each Newton or gradient descent iteration involves sufficient statistics

$$\sum_{i=1}^n \nabla_{\theta}^2 m(W_i, \hat{\theta}_j), \quad \sum_{i=1}^n \nabla_{\theta} m(W_i, \hat{\theta}_j)$$

which can be obtained via distributed computing.

Conclusions

- ▶ We discussed the large p case, which is difficult. Approximate sparsity was used as a generalization of the usual parsimonious approach used in empirical work.
- ▶ A *sea of opportunities* for exciting empirical and theoretical work.
- ▶ We discussed the large n case, which is less difficult. Here the key is the distributed computing. Also big n samples often come in "unlabeled" form, so you need to be creative in order to make good use of them.
- ▶ This is an *ocean of opportunities*.